OSPREYFLOW

Developer Guide

For data engineers: the people who bring data into OspreyFlow. It covers the Ingest Studio from a first pipeline to a scheduled one, the domain model data is mapped into, the rules a load must pass, and the interfaces a data engineer uses beyond the screens. It describes what is built, and says plainly where something is not yet.

Version	1.5
Date	2026-10-28
Describes	OspreyFlow v0.4.0
Status	Issued

Document control

Version 1.5 · 2026-10-28 · For OspreyFlow v0.4.0 · Issued

Version	Date	What changed
1.5	2026-10-28	§6.1: row numbers, counted from 1 in the Glass and from 0 in the API.
1.4	2026-10-27	§7.5: the tests behind it described, not named by path (the build's folder is not part of the product).
1.3	2026-10-27	§7.5: sources withheld for a reader whose case does not read all of an element's datasets, and what a contributor's writer must do about it.
1.2	2026-10-18	§3.3: the NIEM pack's incidents and custody links.
1.1	2026-10-15	A cover page, contents, a list of figures, and this version history.
1.0	2026-10-14	First issue, with screenshots.

Contents

1. How OspreyFlow is built, for a data engineer 5
2. The Data Engineer role 5
3. The Ingest Studio 6
   - 3.1 Pipelines 6
   - 3.2 The canvas 7
   - 3.3 Domain model 8
   - 3.4 Source and connectors 9
   - 3.5 Profile 10
   - 3.6 Map 11
   - 3.7 Load 13
   - 3.8 Verify 14
   - 3.9 Publish, versions and running 14
4. Files 15
5. The ingest folder 15
6. After the load 16
   - 6.1 Held rows 16
   - 6.2 Data feeds 17
   - 6.3 Schedules and watched folders 17
   - 6.4 Reverting a load 17
7. The domain model 18
   - 7.1 Types 18
   - 7.2 What every record carries 18
   - 7.3 Inferred and confirmed 18
   - 7.4 Resolution: never merged 18
   - 7.5 Sources withheld for a scoped reader 18
8. Screening lists 20
9. Interfaces 20
   - 9.1 The API 20
   - 9.2 Files on the server 21
   - 9.3 Commands 21

List of figures

Figure 1 — The Ingest Studio's Pipelines list: name a pipeline, or open one 6
Figure 2 — A pipeline on the canvas: domain model, source, profile, map and load 7
Figure 3 — Choosing the domain pack: the core model, or a pack over it 8
Figure 4 — The Source stage: where the data comes from 9
Figure 5 — The profile: each column's type, gaps, distinct values, samples and pattern matches 10
Figure 6 — The map: source fields on the left, one row's entities and links on the canvas 11
Figure 7 — An entity's settings: its type, what it is matched on, and where each property comes from 12
Figure 8 — The Load stage 13
Figure 9 — Verify: what the recipe would create and reference, and the graph it would make 14
Figure 10 — The Studio's Files tab: images, videos and text files added to the case 15
Figure 11 — Alerts → Held rows 16
Figure 12 — Reports → Data feeds: the pipelines that load into this case, and when each last ran 17

For data engineers: the people who bring data into OspreyFlow. It covers the Ingest Studio from a first pipeline to a scheduled one, the domain model data is mapped into, the rules a load must pass, and the interfaces a data engineer uses beyond the screens. It describes what is built, and says plainly where something is not yet.

The other guides in this folder are the administration guide (administration-guide.pdf), for whoever installs and runs OspreyFlow, and the user guide (user-guide.pdf), for analysts.

---

1. How OspreyFlow is built, for a data engineer

Three tiers; everything that passes between them crosses one of four typed ports:

Tier	What it is
The Glass	The browser workspace. The Ingest Studio is one of its places, shown to anyone holding `edit_pipelines`
The Conduit	The API. It holds the rights, the case levels and the audit, runs the loads, and judges each load with the case's rules
The Vault	The stores: the graph store (ArcadeDB), the analytics database (DuckDB, or ClickHouse), graph compute, the models, Valkey, the object store, OpenSearch, and PostgreSQL for the Conduit's own records

Port	What passes through it
GraphPort	Writing and finding nodes and links, reverting a load's elements, typed traversal
AnalyticsPort	Event aggregates — never raw SQL
ComputePort	Centrality, communities, ownership chains
ModelPort	Extraction, transcription, image description, embedding

Two narrower ports sit beside them: the media store, for content files, and the recognition port.

Access control lives in the stores' queries. Every read is made as a person: their clearance, handling codes and the case's datasets are part of the query itself, never applied afterwards. The Glass never sees a store's own query language or a source system's identifiers.

Three separate questions decide what someone can do with data: their **clearance and handling codes** (what they may see), their **level on a case** — Viewer, Editor or Author (where they may act), and their **rights** (what they may do). A right never widens what a person may see.

2. The Data Engineer role

The role holds	What it lets you do
`edit_pipelines`	Open the Ingest Studio; make and edit pipelines
`verify_pipelines`, `publish_pipelines`, `run_pipelines`, `schedule_pipelines`	Verify, publish, run and schedule them
`add_ingest_files`	List and upload source files
`read_held_rows`, `decide_held_rows`	See, release and reject rows a rule held

The role holds	What it lets you do
`read_data_feeds`, `operate_schedules`	See what feeds a case; Admin → Schedules and runs
`load_screening_lists`, `deploy_rules`, `read_rule_audit`, `read_alert_queue`	Screening lists and the rules that judge loads
`create_case`, `share_case`, `investigate`, `run_analysis`, `read_reports`	Work in cases alongside analysts

It also carries `revert_loads`, for reverting a load — which no screen or command uses yet (section 6.4). It does **not** hold `edit_graph`, `confirm_inferred`, `extract_files`, `resolve_entities`, `export` or `read_report_feed`: those are the Analyst's and Reports Viewer's. An administrator can grant any single right on its own.

3. The Ingest Studio

The Studio is an editor for a **recipe**: a declarative YAML description of how one source becomes part of a case's graph. It works only inside a case — with no case open it says so — and what it loads goes into that case's dataset. A recipe belongs to the first dataset of the case it was made in, and is listed, verified, published and run only in cases that read that dataset.

3.1 Pipelines

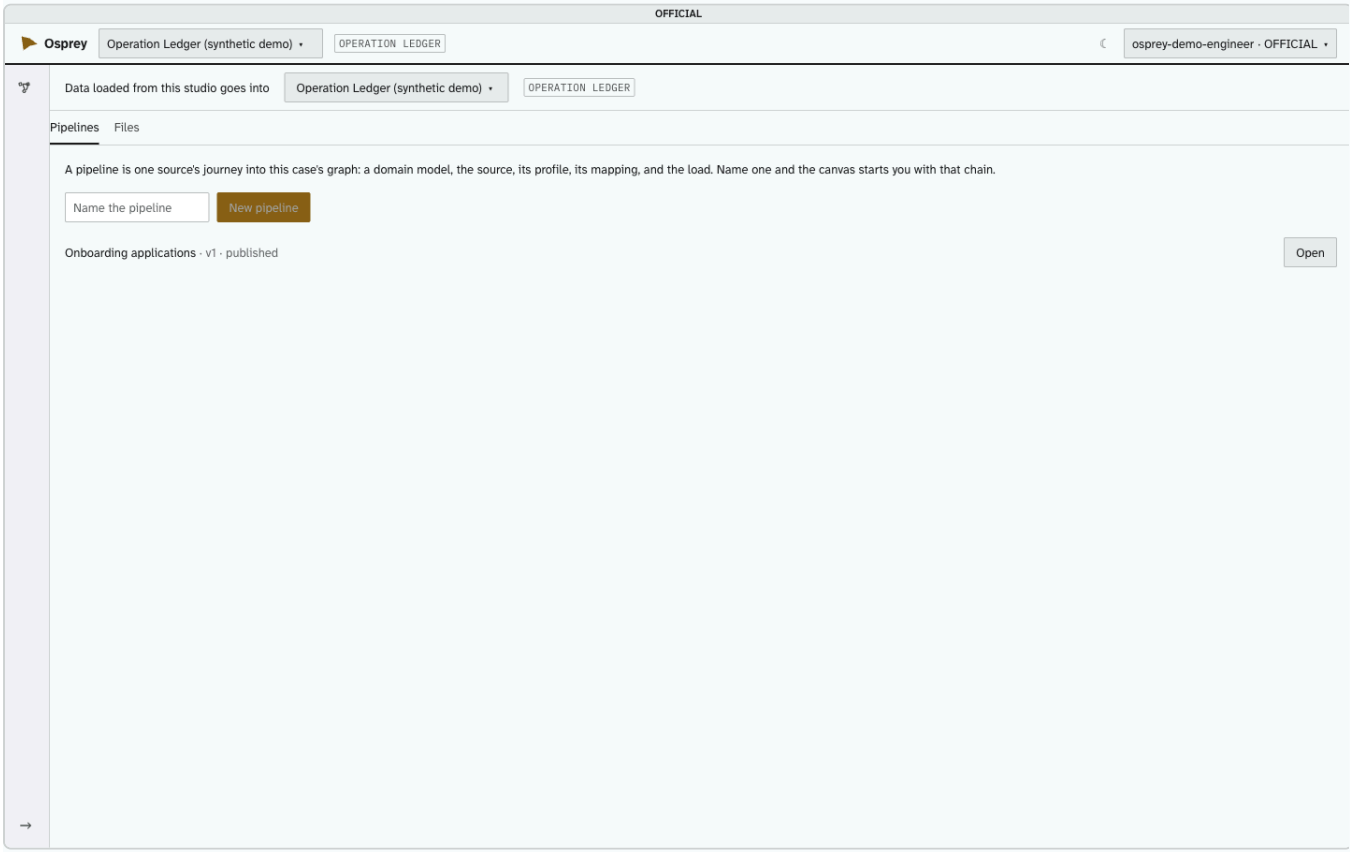


Figure 1. The Ingest Studio's Pipelines list: name a pipeline, or open one

Name the pipeline and choose **New pipeline**: the canvas opens with a whole chain to fill in. Each existing pipeline is listed with its version and status (draft, published or flagged), and **Open** takes you

to it.

3.2 The canvas

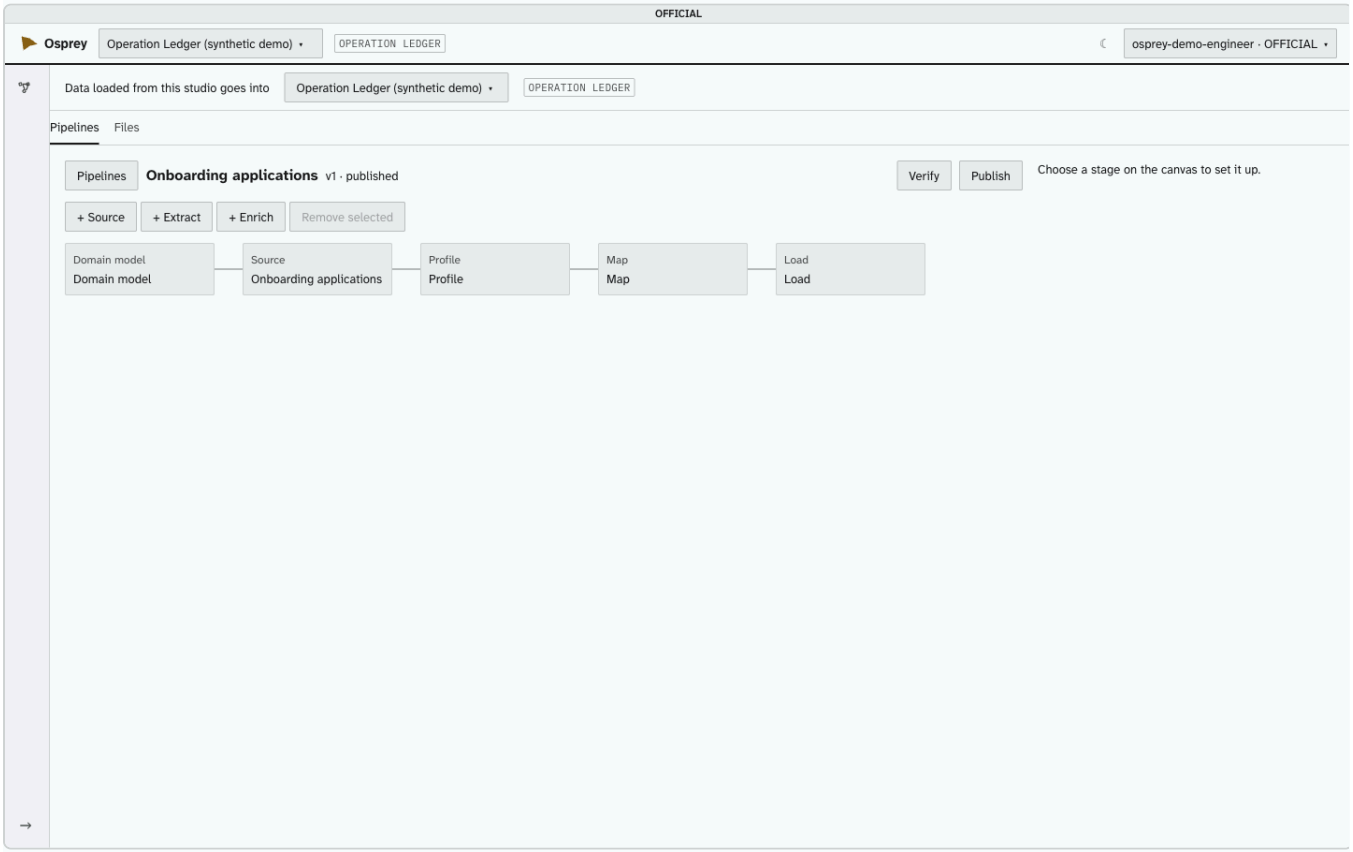


Figure 2. A pipeline on the canvas: domain model, source, profile, map and load

A pipeline is a chain of stages, each with one job:

Stage	What it does
Domain model	What the data means: the core model, and the pack it adds. Exactly one, at the root
Source	The file, its connector, its handling marking and its legal basis
Profile	What the source holds, column by column
Map	How one row becomes entities and links
Extract	Optional: finds entities in the text of Documents a map made
Enrich	Optional: geocodes places — the local gazetteer first
Load	Writes into the case's dataset. It has no settings

- **+ Source** adds another whole source → profile → map → load branch. Each source is its own load, revertible on its own, and inherits the first source's marking and legal basis.
- **+ Extract** and **+ Enrich** go after the last map. **Remove selected** never removes the domain model.
- **Verify** and **Publish** are buttons at the top, not stages.
- Every change is saved as you make it. A pipeline with a cycle, a second domain model, a source that reaches no load, or an extract not fed by a map is refused, in words, before it is saved.

3.3 Domain model

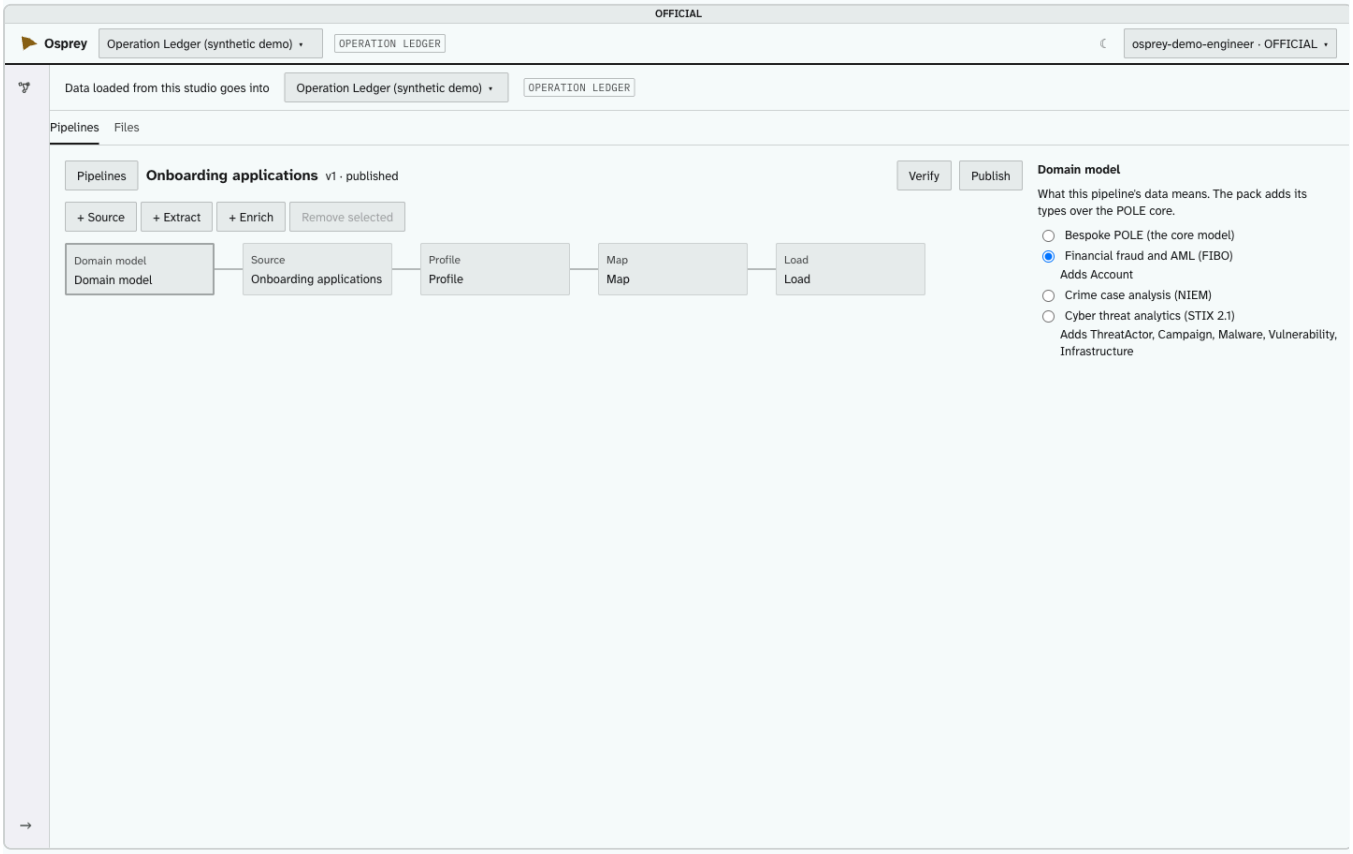


Figure 3. Choosing the domain pack: the core model, or a pack over it

Every pipeline maps into the **POLE core** — people, objects, locations, events, organisations — and may add one pack over it:

Pack	Adds
Bespoke POLE	The core only
Financial fraud and AML (FIBO)	Account; registered at, beneficial owner of (with a percentage and nature), controls, holds account
Crime case analysis (NIEM)	Incident (name, category: assault, disorder, find or other; item; when); used as, role in — each with a required role; shared a cell with (cell, dates), co-defendant of (case reference), in conflict with (reason, dates), visited (in person or video, dates)
Cyber threat analytics (STIX 2.1)	ThreatActor, Campaign, Malware, Vulnerability, Infrastructure; attributed to, targets, indicates, compromises

A pack only adds types, or widens which types a link may join; it cannot change the core.

3.4 Source and connectors

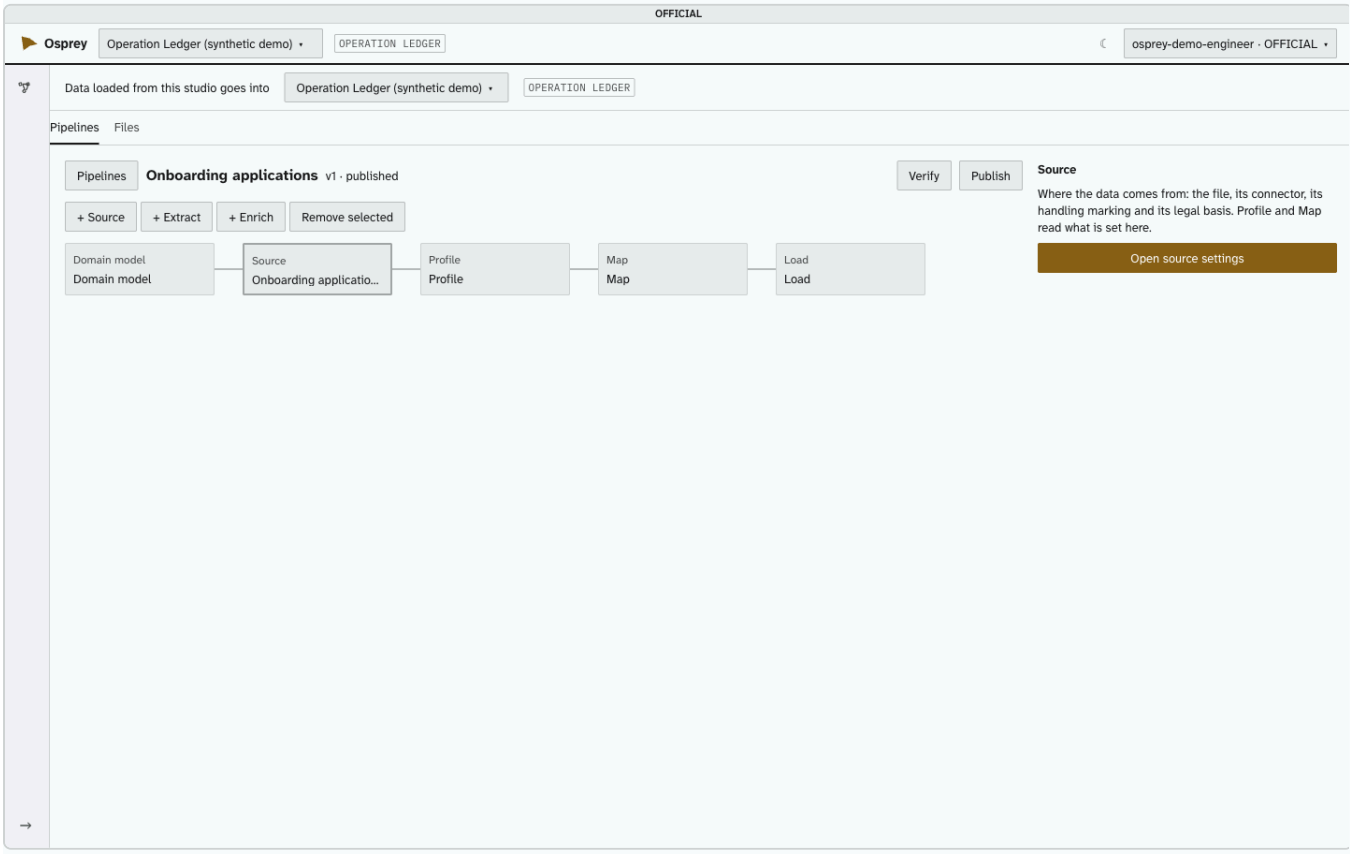


Figure 4. The Source stage: where the data comes from

Open source settings asks three things before anything else can happen:

1. **Where the data is**: a file from the ingest folder, a file from your computer (uploaded into the folder for you), one already on the server, or a path or address you type.
2. **Its handling marking**: the classification and handling code everything from it will carry. A classification left blank is recorded as the scheme's lowest level, and the screen says so.
3. **Its legal basis**: the power or consent the data is held under.

Connectors. The catalogue lists every kind of source OspreyFlow is designed for — communications records, mobile-device extractions, financial records, vehicle and movement data, generic files, documents and media — and marks each as available or not. **Available today: four connectors, all reading CSV**: generic CSV, call records (day-first dates), cell sites (WGS 84 co-ordinates) and seizure records. The others are shown as not yet available. A connector is a small profile file in `connectors/`; adding one needs the Conduit restarted.

3.5 Profile

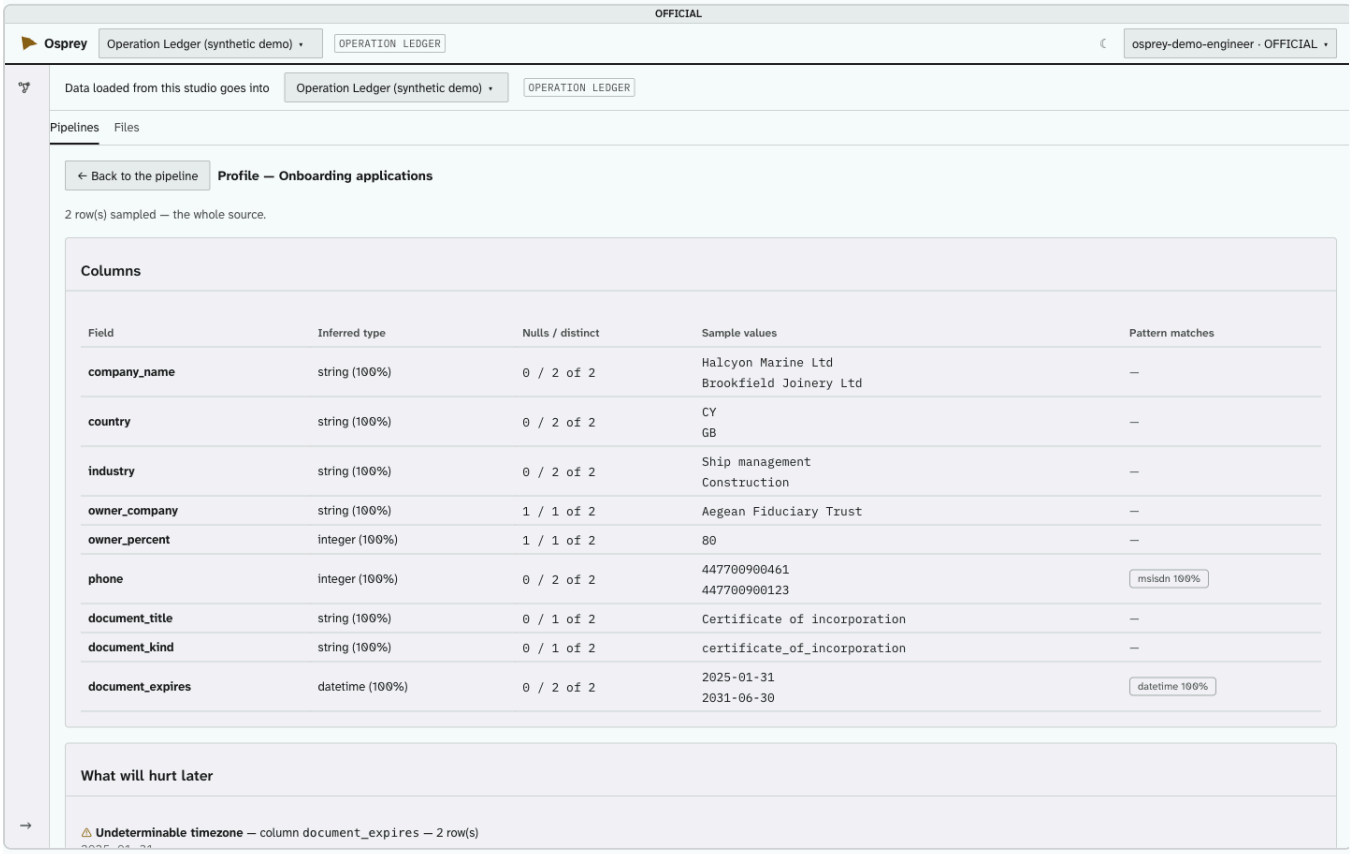


Figure 5. The profile: each column's type, gaps, distinct values, samples and pattern matches

The profile reads up to 10,000 rows (or the whole source, if smaller) and shows, for each column, its inferred type and how sure that is, how often it is empty, how many distinct values it has, five sample values, and which known patterns it matches (a phone number, an IMEI, a date…) and how often. A share is never rounded up to 100% when it is not.

What will hurt later names the problems before they reach the graph: mixed date formats, a timezone that cannot be told, an encoding other than UTF-8, rows with differing column counts, and apparent duplicates.

Someone not cleared for the source's marking sees the samples' shape only: letters as `A`, digits as `#`. Map cannot be opened until the profile has run.

3.6 Map

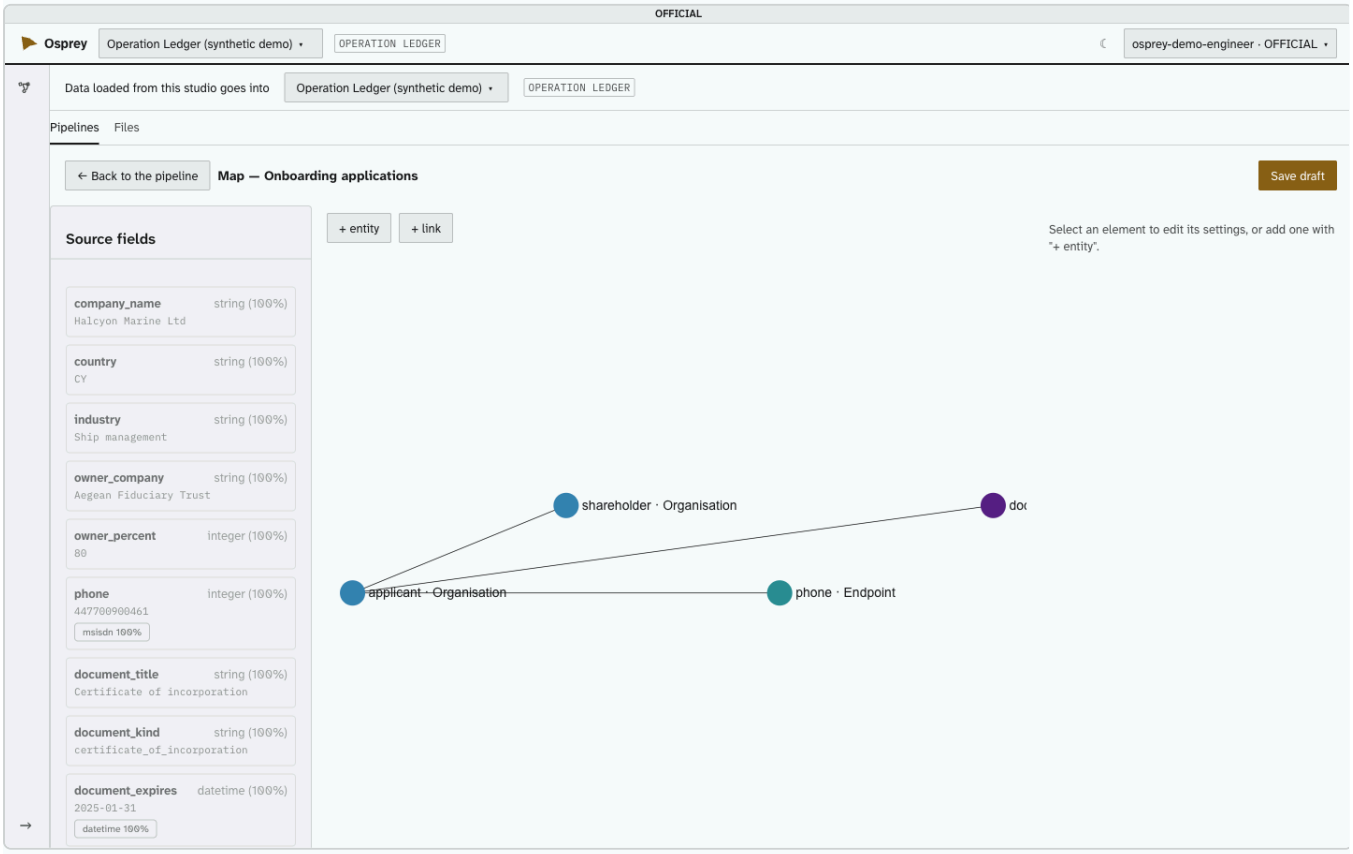


Figure 6. The map: source fields on the left, one row's entities and links on the canvas

One row of the source becomes a small graph. **+ entity** adds an entity — first its POLE class, then its type; **+ link** joins two. The source's fields are listed on the left with a sample of each; every field not yet used is shown until it is mapped or marked **ignore**.

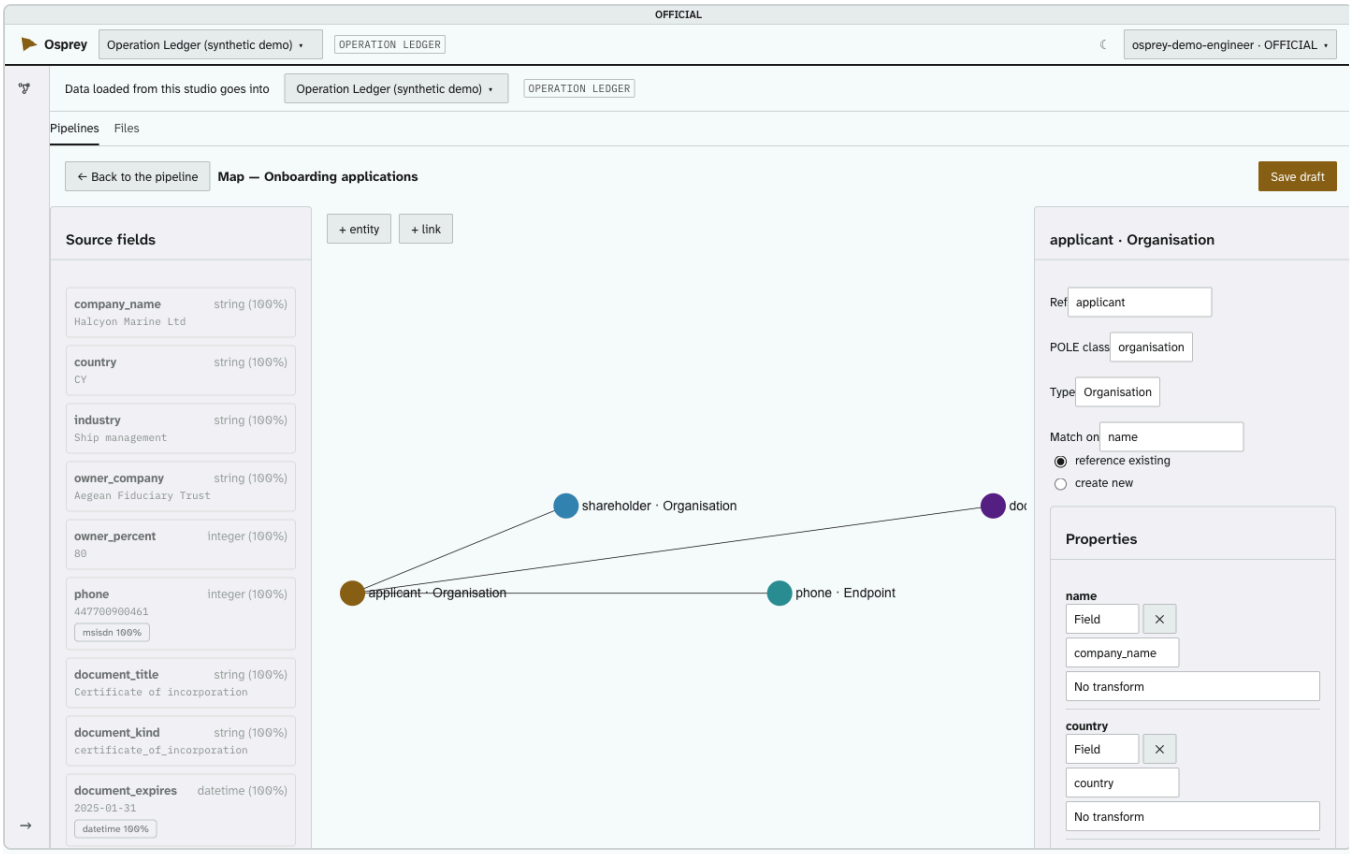


Figure 7. An entity's settings: its type, what it is matched on, and where each property comes from

For each entity:

- **Match on** — the property that identifies it — and whether a row **references an existing** entity with that value or **creates a new** one.
- **Properties**: each takes a field, several fields, a fixed value, or a value map, optionally through a transform.
- **When**: an element can depend on a field's value.
- **Content this holds**: a Document's text, an image or a video.

For each link: its type, and its time — not recorded, or taken from a column. A link with a time can mark text a model may later read.

Transforms are Knowledge Adapters, never expressions. OspreyFlow has no expression or scripting language. A transform is one of a fixed set of tested adapters: `msisdn`, `imei`, `imsi`, `email`, `datetime`, `person_name`, `address`, `iban`, `card`, `coordinate_lon`, `coordinate_lat`, `geometry_ref`, `polygon_ref`, `grid_ref` and `postcode`. Some take more than one column (co-ordinates two, a grid reference three). A structured result is bound one part at a time, such as `person_name(part=family)`. A transform that does not exist yet is a new adapter, written with golden test cases — not something typed into a recipe.

Handling and provenance default from the source; an element may **tighten** its handling, never loosen it. Data a recipe loads is *asserted* — it came from a source, not a model.

Refused at save, in words: a link of a type the ontology marks as probabilistic, `SAME_AS` or `NOT_SAME_AS` (resolution writes those, never a recipe), an unknown type, a property the ontology does not have, and a link's model-readable text with no time.

Save draft saves the map; it does not verify it.

3.7 Load

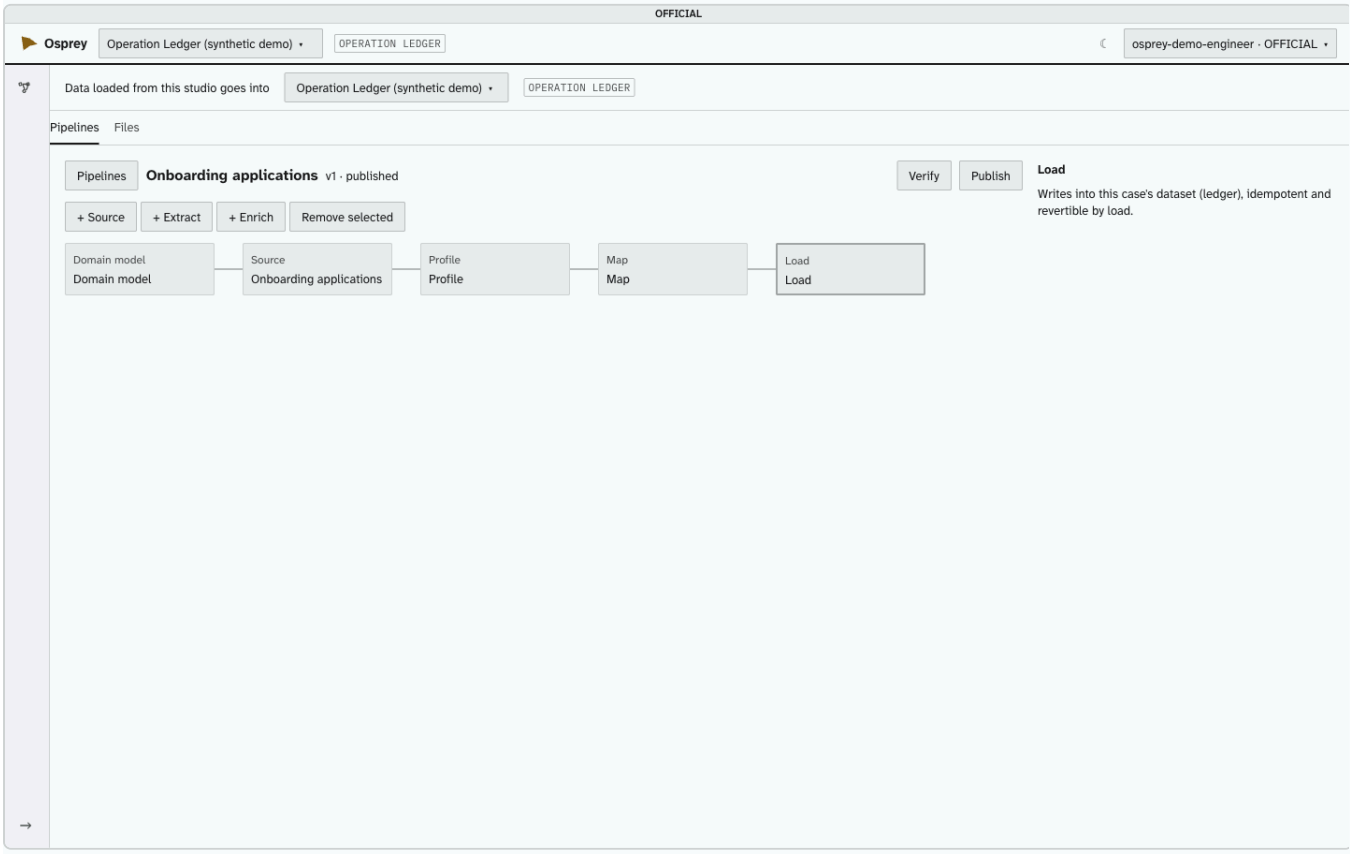


Figure 8. The Load stage

The load writes into the case's dataset. It is **idempotent** — running the same data again adds nothing new — and every element it writes carries its load's id, so a load can be reverted as a whole.

3.8 Verify

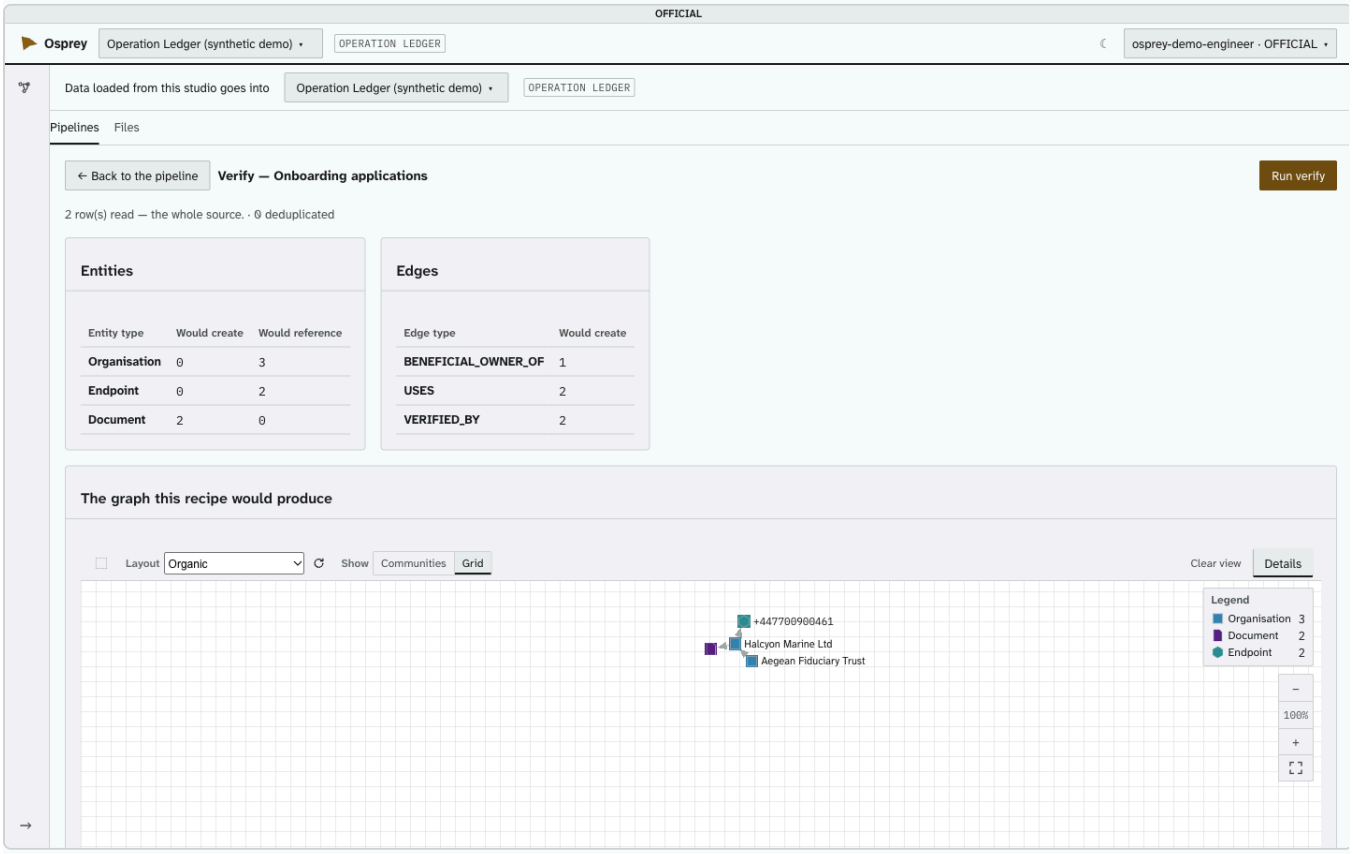


Figure 9. Verify: what the recipe would create and reference, and the graph it would make

Verify is a dry run: nothing is written. It shows, for every type, how many entities and links the recipe **would create**, how many it **would reference** (existing ones, looked up within the case's datasets as you), how many rows were deduplicated or rejected, and the graph itself.

Rejections are grouped by cause, naming the adapter and field — never the value. A value the ontology refuses (a date of birth that is not a date) rejects its row; nothing is silently changed. **A recipe is published only after it has verified.**

3.9 Publish, versions and running

Publish commits the recipe to `recipes/<id>.yaml` in the deployment's recipe repository and marks it published. To change a published recipe, **Reopen for editing** makes a new draft version; a published version is never edited in place. If the ontology changes under a published recipe, the recipe is **flagged**, and its scheduled runs stop until it is looked at again.

A publish is refused if a stage cannot run here: an extract library this deployment does not have, an unknown geocoder, or a handling override with no reason.

Run now loads the source. Every row is staged in the object store first, then the graph is written, then events go to the analytics database and places to the map. The run's report gives rows read, rows deduplicated, entities created and referenced, links created, rejections, rows held, and what the case's rules said. A run is refused if the source is marked above your clearance.

Known limits: a re-run adds its events to the analytics database a second time (the graph is unaffected); and loads write one element at a time, so a large source takes a while.

4. Files

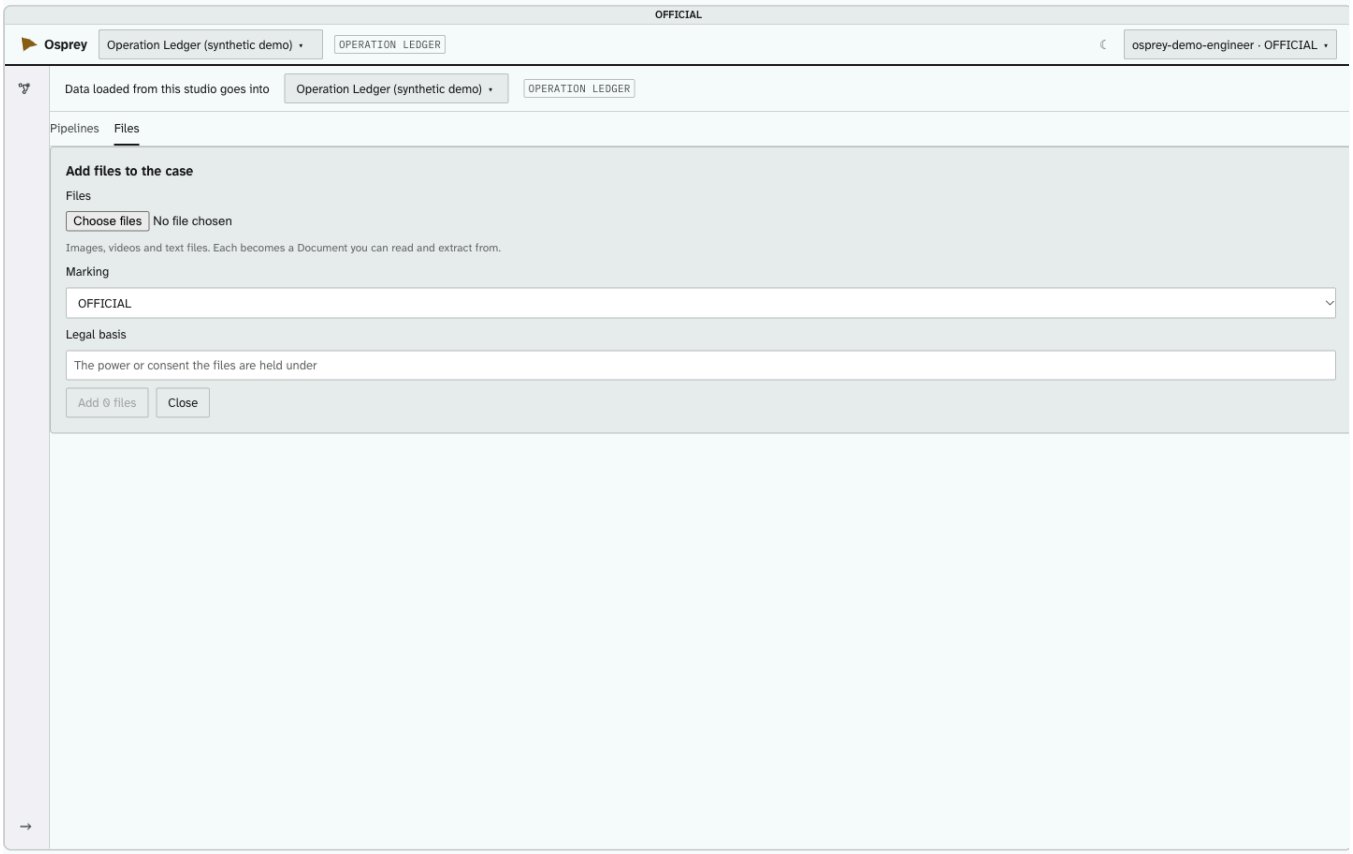


Figure 10. The Studio's Files tab: images, videos and text files added to the case

The Studio's **Files** tab adds images, videos and text files to the case, each becoming a Document that can be read and extracted from (PNG, JPEG, GIF, WebP, TIFF, HEIC, MP4, MOV, WebM and plain text). It takes a marking and a legal basis for the files. PDF and Word documents are not read yet. **Adding files there needs `extract_files`**, which the Data Engineer role does not hold — an analyst adds them, or an administrator grants the right.

5. The ingest folder

The source files a pipeline reads live in the deployment's **ingest folder** (`components.storage.ingest`). OspreyFlow only ever **reads** it, except for its `uploaded/` folder, where files uploaded through the Studio land, under the case's dataset. The Studio lists files directly in the folder, in `uploaded/`, and in each of the case's dataset folders — never anything outside the folder, in deeper folders, or hidden.

An upload is refused, in words, if its name has a folder in it, starts with a dot or is over 200 characters; if it is empty, or over the size limit (512 MiB by default); if its kind is not one OspreyFlow takes; or if a file of that name is already there — **nothing is ever overwritten**.

6. After the load

6.1 Held rows

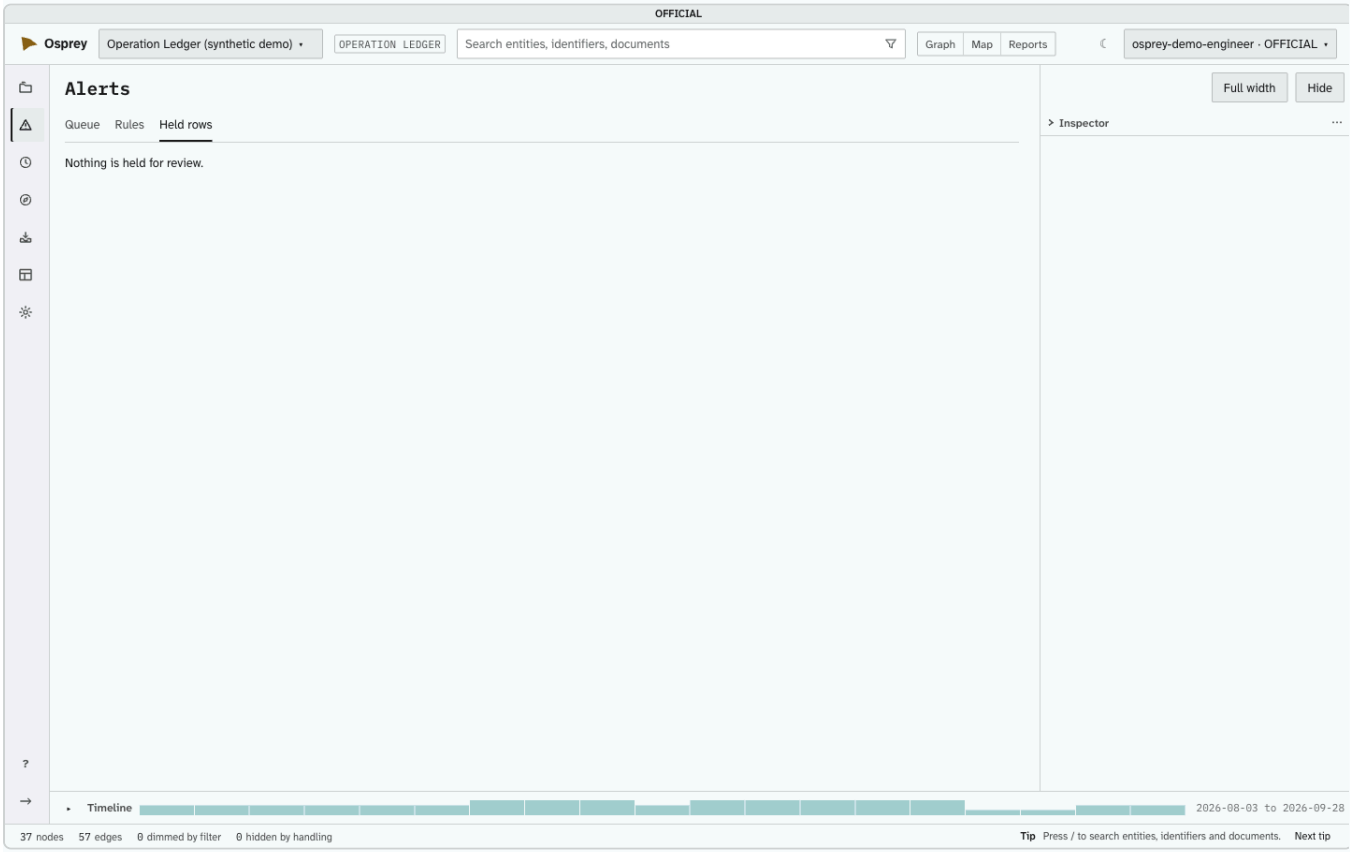


Figure 11. Alerts → Held rows

A case's rules judge every load before it commits. A row a rule could not judge is **held**: it reaches no store until someone decides it, in the Glass under **Alerts → Held rows**, which shows why each row was held. **Release and load** writes it under a new load (it is judged again, and never held a second time); **Reject** needs a written reason. Listing held rows needs Editor on the case; **deciding them needs Author**.

Row numbers. The Glass counts rows from 1: the row under the header is row 1, in held rows, Verify's rejected rows and a source's `row N`. The API and what is stored count from 0: a `row N` locator, a rejection's or held row's `row_index`, and an alert's `fired_on` `row`. A tool of your own that shows one of these to a person adds one. Sample rows (`…/map/sample-rows/{row_index}`) are indexed from 0 too.

6.2 Data feeds

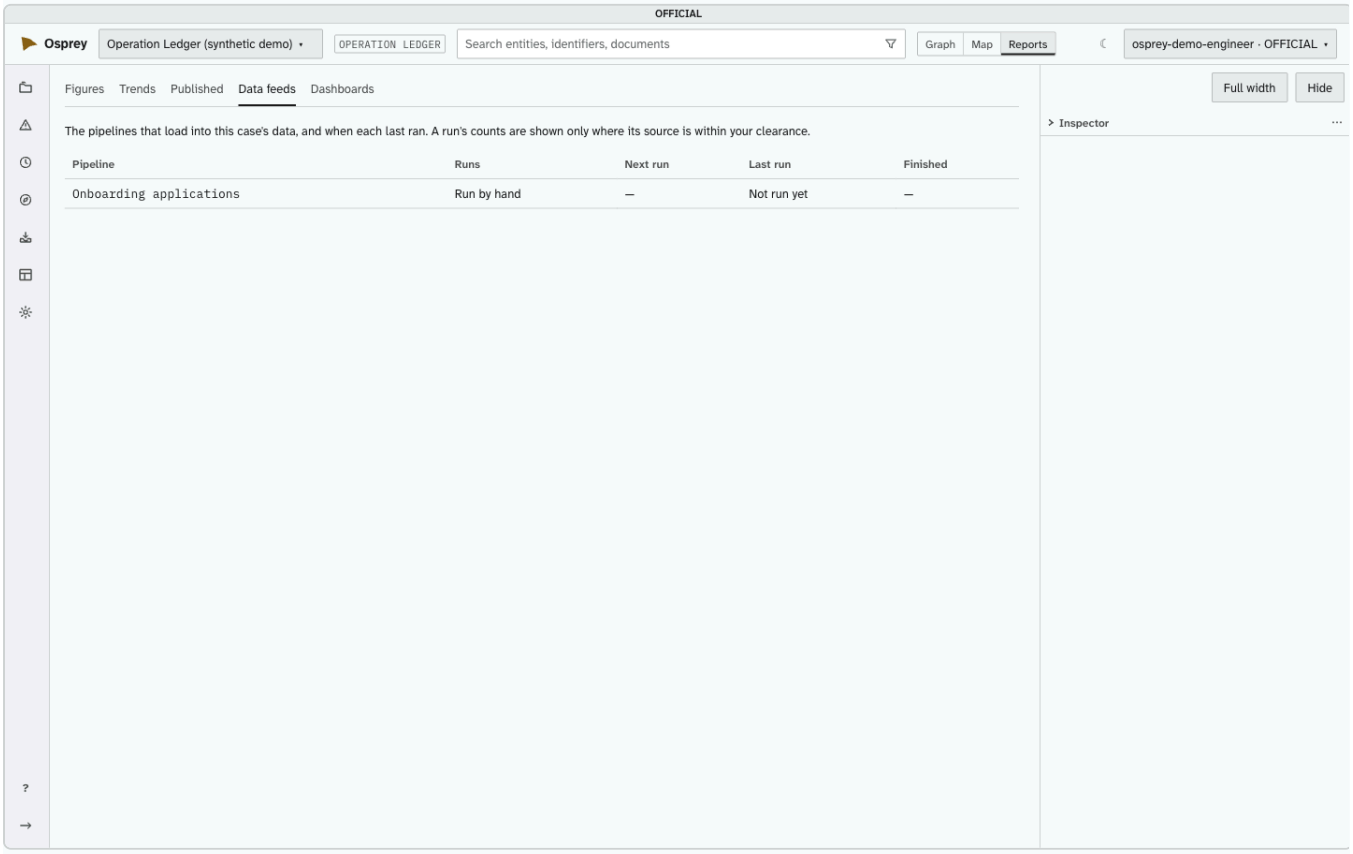


Figure 12. Reports → Data feeds: the pipelines that load into this case, and when each last ran

Under **Reports → Data feeds**, anyone holding `read_data_feeds` with Viewer on a case sees the pipelines that load into it, how they run (by hand or on a schedule), and when each last ran. A run's counts are shown only to someone cleared for its source.

6.3 Schedules and watched folders

A published pipeline can run on a **schedule** — every hour, every day, every weekday, every week, or a cron expression — in a named time zone (Europe/London by default), never more often than every 15 minutes, since each run is a whole load. The screen previews the schedule in words with its next five runs.

It can also run when a **file arrives** in its dataset's own folder, matching a pattern such as `*.csv`: each file is taken once, a minute after it stops changing, one at a time.

Either way it runs **as the person who last set it**, with every check their own run would get. A run that is refused — they can no longer sign in, have lost the right, are no longer an Editor, or are cleared below the source — tells them on their alerts; **three refusals in a row pause it**. Admin → Schedules and runs lists every one.

6.4 Reverting a load

Every element a load writes carries its load's id, and reverting a load hides all of them from every read without deleting anything. **The revert exists in the product's code but is not yet reachable from a screen or a command**; when it is, it will need `revert_loads`.

7. The domain model

7.1 Types

The ontology (`packages/ontology/default.yaml`) defines the entity types — **Person, Endpoint, Device, Location, Organisation, Event, Document, Vehicle, Asset** — and the link types: communicated with, uses, present at, mentioned in, verified by, holds, depicts, associate of, family of, member of, transacted, paired with. A deployment may add types; it cannot add POLE classes. An Endpoint's value is always held normalised, with its kind beside it.

The graph holds **one summarised link per pair per type**, with its first and last time, how often it occurred and its weight; the individual events are in the analytics database.

7.2 What every record carries

Part	What it holds
Provenance	Asserted, inferred or unresolved; its sources (withheld from a reader whose case does not read all of the element's datasets, section 7.5); what extracted it and how sure; when first seen, last confirmed, and by whom
Handling	Classification, handling code, action code, sanitisation code, originator, owning unit, review date, legal basis
Datasets	Which datasets it belongs to; a later load adds, never removes
Audit	The load, recipe and recipe version that wrote it; its revision; whether it was reverted

Handling never goes down. Writing an element again never lowers its marking. A clearance the deployment does not recognise sees nothing.

7.3 Inferred and confirmed

Anything a model produces is **inferred** until a person confirms it; confirming records who and when. Nothing is promoted automatically, and no confidence score does it on anyone's behalf.

7.4 Resolution: never merged

Loads write each source's own records, and **nothing is ever merged in storage**. Resolution links records it thinks are one with a scored `SAME_AS`, and the Conduit shows each linked group as one at the moment it is read, at the case's chosen strictness: Strict (0.99), Balanced (0.95) or Broad (0.80). An analyst can link two records by hand, or keep two apart — which is always honoured. A recipe may not write either.

7.5 Sources withheld for a scoped reader

An element can belong to more than one dataset, and a case reads only the datasets it names. A reader may see an element through any one of its datasets, but its **sources** (the files and rows it came from) name no dataset. So the sources are sent to a reader only when every dataset the element belongs to is one the reader's case reads. Otherwise the reader gets the node or link with its sources emptied and a flag set: what it came from is held back, not the fact that it exists.

Where it is applied. In the Vault, where an element is read for a principal, so that nothing stored or forwarded carries the sources past it:

- The graph adapter's element mapping, `node_from_element_map` and `edge_from_element_map` in `packages/adapters/graph_arcadedb/serialization.py`, which every read path goes through, given the reading principal.
- The working set's projection, `_project` and `project_delta` in `services/api/working_set/service.py`, which serve the replay, the live stream and a restore.
- A principal whose `datasets` is `None` is the Conduit's own unlimited work, and reads every dataset (`services/api/dataset_scope.py`). Nothing a request handler is given is that principal.

The shape. `Provenance.sources_withheld` (`packages/schemas/domain.py`) is `False` by default. When it is `True`, `sources` is empty; when it is `False`, `sources` is never empty. The generated TypeScript types follow. The predicate is one function, `sources_within_case`, in `packages/ontology/handling.py`: every dataset of the element is in the principal's `datasets`; an element that names no dataset, read by a scoped principal, is withheld. `withhold_sources` applies it and returns a copy; it changes nothing stored.

The rule for a writer that reads and writes back. A writer that reads an element and writes its provenance back (a link change, a picture, an extraction's replacement, a resolution, a pipeline run's mention, confirming an alert's cited links) would rebuild the element from a withheld, empty list. That would silently drop the stored sources, or be refused by the write guard below. So **a read whose result is written back runs inside `ontology.handling.sources_in_full()`**:

```
from ontology.handling import sources_in_full

with sources_in_full():
    found = await graph.get_edges([edge_id], principal)
# change what you mean to change in found[0], then write it back
```

The read is still made with the request's principal, so the access filter is applied in the adapter as before; only the sources are left in. The scope is a context variable, so it holds for the calling task alone, and the adapter and its test double honour it. The writers that do this today are `services/api/link_changes.py`, `services/api/entity_pictures.py`, `services/extract/writer.py`, `services/resolve/trigger.py`, `services/api/pipeline_run.py` and `services/api/alert_queue.py`. The working set's own writes are covered the same way. Never wrap a read whose result goes to the Glass.

What such a writer sends back. Whatever a writer returns or streams to the Glass (a delta's `to_sent()`, a route's response) passes through `withhold_sources` for the reader first. A writer that read in full has the sources in its hands, and a response built from them would send the sources the reader may not have. `services/api/link_changes.py` does this for the edge it returns, and `services/api/working_set/models.py` for the deltas it sends.

A withheld provenance is never written. The write guard, `_refuse_withheld` in `packages/adapters/graph_arcadedb/serialization.py`, raises if a `NodeWrite` or `EdgeWrite` carries a `Provenance` with `sources_withheld` set, since writing it would replace the stored sources with none. A refusal there means a read that was written back was not inside `sources_in_full()`.

For a new writer. Ask two questions. Does it read an element and write its provenance back? Then wrap the read. Does it send the element, or anything built from it, to a reader? Then pass it through `withhold_sources` first. The build's tests hold every writer to this: each reads as a scoped principal, changes an element in two datasets, and checks that every stored source is kept and that what is sent back carries none.

8. Screening lists

Sanctions, PEP and adverse-media lists are kept beside the graph, not in it, deployment-wide, and replaced a whole list at a time. They are loaded from a YAML file on the server:

```
python -m scripts.load_screening_list <file.yaml>
```

The file names the list, whether it is synthetic, and its entries — each with an id, a listing (`sanction`, `pep` or `adverse_media`), an entity type (Person or Organisation), a name, and optionally a date of birth and country. **A list of real people is refused** unless the deployment's local real-data exceptions name its dataset. A match reaches the graph only as a scored, inferred `SAME_AS`, for a person to confirm.

9. Interfaces

9.1 The API

Every route is under `/v1`; a case's routes are under `/v1/investigations/{investigation_id}`. A missing right is refused with `403` and the right's name; a case you have no level on answers as if it did not exist.

Method and path	Needs
`GET /v1/ingest/connectors`, `/v1/ingest/packs`, `/v1/ingest/transforms`	Signed in
`GET`, `POST …/ingest/source-files`	`add_ingest_files` (upload: Editor)
`POST …/ingest/recipes`; `PATCH`, `GET …/ingest/recipes/{recipe_id}`; `GET …/ingest/recipes`	`edit_pipelines`
`PUT …/ingest/recipes/{recipe_id}/pipeline`	`edit_pipelines`, Editor
`GET …/ingest/recipes/{recipe_id}/profile`	`edit_pipelines`
`PUT …/row-template`; `GET …/map/sample-rows/{row_index}`; `POST …/map/handling-preview`	`edit_pipelines`
`POST …/ingest/recipes/{recipe_id}/verify`	`verify_pipelines`, Editor
`POST …/ingest/recipes/{recipe_id}/publish`; `POST …/versions`	`publish_pipelines` / `edit_pipelines`, Editor
`POST …/ingest/recipes/{recipe_id}/run`	`run_pipelines`, Editor
`PATCH …/schedule`; `POST …/schedule/pause`, `…/resume`; `POST /v1/ingest/schedule-preview`	`schedule_pipelines`
`GET …/ingest/recipes/{recipe_id}/runs`; `GET …/data-feeds`	`read_data_feeds`, Viewer
`GET …/held-rows`	`read_held_rows`, Editor
`POST …/ingest/recipes/{recipe_id}/held-rows/release`; `POST …/held-rows/reject`	`decide_held_rows`, Author
`GET /v1/admin/schedules` and its actions	`operate_schedules`
`GET /v1/me/rights`	Signed in: your own rights

9.2 Files on the server

Where	What
The ingest folder (`components.storage.ingest`)	Source files, read-only to OspreyFlow; `uploaded/` beneath it is where uploads land
`recipes/<id>.yaml`	Each published recipe
`connectors/<id>.yaml`	Each available connector's profile
`packages/ontology/default.yaml`, `packs/`	The ontology and its domain packs

9.3 Commands

Command	What it does
`ospreyflow demo list`, `ospreyflow demo add <id>`	The synthetic demo packs, each loaded into its own case
`python -m scripts.seed_ingest_folder`	Writes the demo packs' synthetic source files into the ingest folder
`python -m scripts.load_screening_list <file.yaml>`	Loads or replaces a screening list (section 8)